



GCP REIMAGINE IDEATHON 2023

gcdc-oht33765u9-2023

TEAM MEMBERS

Mayank Harlalka

(Analyst – Data Science)

DNA (ANALYTICS)

CHD_SEZ









Mayank.harlalka@infosys.com

+919988003987



SOLUTION : 3 USE CASES

1. Trap Camera Animal Detection
2. Drone Search and Rescue
3. Driver Drowsiness Detection

<div><div> Filter</div><div>Filter buckets</div></div> <div><div></div><div></div></div>							
<input type="checkbox"/>	Name 	Created	Location type	Location	Default storage class 	Last modified	Public ac
<input type="checkbox"/>	drone-search-rescue-data	Jul 8, 2023, 3:00:28 PM	Region	europe-west6	Standard	Jul 8, 2023, 3:00:28 PM	Not pub 
<input type="checkbox"/>	drowsiness-data	Jul 5, 2023, 11:38:19 PM	Region	us-central1	Standard	Jul 5, 2023, 11:38:19 PM	Not pub 
<input type="checkbox"/>	trap-camera-data	Jul 5, 2023, 10:27:33 AM	Region	europe-west9	Standard	Jul 20, 2023, 12:15:55 PM	Not pub 

TRAP-CAMERA ANIMAL DETECTION

UseCase Type:

Model Deployment : Cloud

Detection : Cloud

Storage : Cloud

Action Taken : On Cloud (Alerting Concerned Services/ People)

OVERVIEW



Forest Rangers, Wildlife Conservationists etc. setup Trap Cameras in the wild to monitor the activities of animals and to get photographs of animals that tend to be elusive.

Hours are spent manually going through each and every video to across even a tiniest hint of an animal.

Hence we can automate this process of going through the videos and get the animals detected easily.



CURRENT WORKING ARCHITECTURE



Videos from Tap Camera come in batches daily.

e.g. 24 1-hour videos for yesterday

The videos land in a particular folder inside a bucket in Cloud Storage

Vertex AI jupyter notebooks have pre-defined python code, that can be run to start the detection in videos

The code automatically stores the detections in the bucket. The detections are images with bounding boxes around the animal found. Also, the detections are named as VideoName-Timestamp.jpg

As soon as new detections are stored in the bucket, Cloud PubSub will trigger a notification to subscribed web services, which can ingest the topic to send email/SMS to concerned people

EXAMPLE SCENARIO



In our Wildlife Park in Indonesia, we have the following 5 species of concern, which we would like to track and monitor:

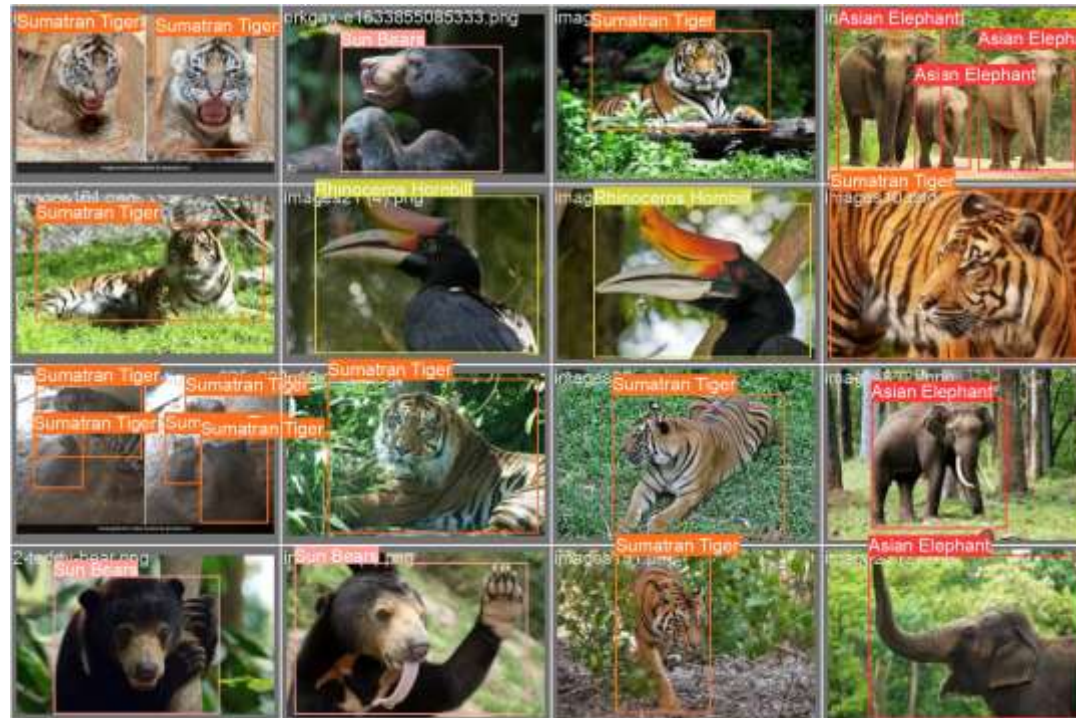
- Sumatran Tiger (*Panthera tigris sumatrae*)
- Asian Elephant (*Elephas maximus*)
- Pig Tailed Macaque (*Macaca nemestrina*)
- Rhinoceros Hornbill (*Buceros rhinoceros* Linnaeus)
- Sun Bears (*Helarctos malayanus*)

We would like to get notified about the image capture of any of these in the previous day. The images should contain the detection bounding box across the animal, with its name.

Also, the image file name should contain the video name, and the timestamp of detection in the video

MODEL USED

YOLOv8 Object Detection Model (<https://docs.ultralytics.com/>)



MODEL TRAINING

The model was trained on Vertex AI Notebook, with 800+ images collected and labelled manually across the 5 classes (animals), in the YOLO format.

The dataset was further augmented to include multiple scenarios for light, color, distortions etc.

```
Epoch  GPU_mem  box_loss  cls_loss  dfl_loss  Instances  Size
199/200   13.4G    0.2841   0.2162   0.8907      2         640: 100%|██████████| 42/42 [00:17<00:00, 2.36it/s]
          Class   Images  Instances  Box(P      R      mAP50  mAP50-95): 100%|██████████| 6/6 [00:02<00:00, 2.78it/s]
          all     165      207    0.929    0.851    0.905    0.742

Epoch  GPU_mem  box_loss  cls_loss  dfl_loss  Instances  Size
200/200   13.5G    0.2767   0.2186   0.8881      2         640: 100%|██████████| 42/42 [00:17<00:00, 2.45it/s]
          Class   Images  Instances  Box(P      R      mAP50  mAP50-95): 100%|██████████| 6/6 [00:04<00:00, 1.45it/s]
          all     165      207    0.918    0.861    0.905    0.742

200 epochs completed in 1.786 hours.
Optimizer stripped from tmp/model/runs_0/train2/weights/last.pt, 136.7MB
Optimizer stripped from tmp/model/runs_0/train2/weights/best.pt, 136.7MB


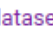

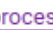
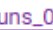
Validating tmp/model/runs_0/train2/weights/best.pt...
Ultralytics YOLOv8.0.136 Python-3.10.11 torch-2.0.1+cu117 CUDA:0 (Tesla V100-SXM2-16GB, 16161MiB)
Model summary (fused): 268 layers, 68128383 parameters, 0 gradients
          Class   Images  Instances  Box(P      R      mAP50  mAP50-95): 100%|██████████| 6/6 [00:03<00:00, 1.57it/s]
          all     165      207    0.928    0.844    0.91    0.745
Asian Elephant    165      46    0.959    0.87    0.958    0.739
Sun Bears         165      35    0.936    0.857    0.925    0.776
Sumatran Tiger    165      63    0.898    0.714    0.796    0.652
Pig Tailed Macaque 165      33    0.883    0.909    0.917    0.722
Rhinoceros Hornbill 165      30    0.963    0.868    0.952    0.835

Speed: 0.4ms preprocess, 8.0ms inference, 0.0ms loss, 5.0ms postprocess per image
Results saved to tmp/model/runs_0/train2
```

Buckets > trap-camera-data > model

UPLOAD FILES UPLOAD FOLDER CREATE FOLDER

Filter by name prefix only Filter Filter objects and folders

<input type="checkbox"/>	Name	Size
<input type="checkbox"/>	 bestn.pt ← Final Model	5.9 MB
<input type="checkbox"/>	 dataset_0/	—
<input type="checkbox"/>	 logs/	—
<input type="checkbox"/>	 process_video.py	3.7 KB
<input type="checkbox"/>	 runs_0/	—

```
YOLO Model Train.ipynb
+
Python 3
•[13]: import os
      from google.cloud import storage
      # Initialize Google Cloud Storage client
      storage_client = storage.Client()
      # Specify your bucket name and model folder
      bucket_name = "trap-camera-data"
      model_folder_name = "model"
      # Define the temporary local directory path within the Vertex environment
      temp_model_dir = "/tmp/model/"
      # os.makedirs(temp_model_dir)
      # Get the list of blobs (files) in the model folder
      bucket = storage_client.bucket(bucket_name)
      blobs = bucket.list_blobs(prefix=model_folder_name + "/") # Specify the model folder as the prefix
      # Copy each blob (file) from the model folder to the temporary local directory
      for blob in blobs:
          if not blob.name.endswith('/'): # Check if the blob is a file and not a directory
              # Construct the destination path for the file in the temporary directory
              relative_path = os.path.relpath(blob.name, model_folder_name) # Get the relative path within the model folder
              destination_path = os.path.join(temp_model_dir, relative_path)
              # Create subdirectories if needed
              os.makedirs(os.path.dirname(destination_path), exist_ok=True)
              # Download the file to the destination path
              blob.download_to_filename(destination_path)
              print(f"Copied {blob.name} to {destination_path}")
      # Now you can access the model files in the temporary local directory (e.g., /tmp/model)

Copied model/bestn.pt to /tmp/model/bestn.pt ***

[15]: model.train(data=r"/tmp/model/dataset_0/dataset.yaml", epochs=200, project=r"/tmp/model/runs_0" )

Ultralytics YOLOv8.0.136 Python-3.10.11 torch-2.0.1+cu117 CUDA:0 (Tesla V100-SXM2-16GB, 16161MiB)
```

SAMPLE OPERATION — STEP 1

As an example, 5 videos were placed in a specific folder:

[←](#) Bucket details [REFRESH](#) [HELP ASSISTANT](#) [LEARN](#)

trap-camera-data

Location	Storage class	Public access	Protection
europe-west9 (Paris)	Standard	Not public	None

[OBJECTS](#) [CONFIGURATION](#) [PERMISSIONS](#) [PROTECTION](#) [LIFECYCLE](#) [OBSERVABILITY](#) [INVENTORY REPORTS](#)

Buckets > trap-camera-data > new_data

[UPLOAD FILES](#) [UPLOAD FOLDER](#) [CREATE FOLDER](#) [TRANSFER DATA](#) [MANAGE HOLDS](#) [DOWNLOAD](#) [DELETE](#)

Filter by name prefix only Filter | Filter objects and folders Show deleted data

<input type="checkbox"/>	Name	Size	Type	Created	Storage class	Last modified	Public access	
<input type="checkbox"/>	1.mp4	2.6 MB	video/mp4	Jul 20, 2023, 3:29:25 PM	Standard	Jul 20, 2023, 3:29:25 PM	Not public	
<input type="checkbox"/>	2.mp4	5.2 MB	video/mp4	Jul 20, 2023, 3:29:41 PM	Standard	Jul 20, 2023, 3:29:41 PM	Not public	
<input type="checkbox"/>	3.mp4	36.4 MB	video/mp4	Jul 20, 2023, 4:06:17 PM	Standard	Jul 20, 2023, 4:06:17 PM	Not public	
<input type="checkbox"/>	4.mp4	5.7 MB	video/mp4	Jul 20, 2023, 3:58:10 PM	Standard	Jul 20, 2023, 3:58:10 PM	Not public	
<input type="checkbox"/>	5.mp4	7.4 MB	video/mp4	Jul 20, 2023, 4:06:31 PM	Standard	Jul 20, 2023, 4:06:31 PM	Not public	

SAMPLE OPERATION – STEP 2

After the model was run in the Vertex AI notebook, it created the detection images in separate folder for each video, with the image name as VideoName-Timestamp.jpg

Buckets > trap-camera-data > detections

UPLOAD FILES UPLOAD FOLDER CREATE FOLDER TRANSFER DATA

Filter by name prefix only Filter Filter objects and folders

<input type="checkbox"/>	Name	Size	Type	Created	Storage
<input type="checkbox"/>	1/	—	Folder	—	—
<input type="checkbox"/>	2/	—	Folder	—	—
<input type="checkbox"/>	3/	—	Folder	—	—
<input type="checkbox"/>	4/	—	Folder	—	—
<input type="checkbox"/>	5/	—	Folder	—	—

Buckets > trap-camera-data > detections > 3

UPLOAD FILES UPLOAD FOLDER CREATE FOLDER TRANSFER DATA

Filter by name prefix only Filter Filter objects and folders

<input type="checkbox"/>	Name	Size	Type
<input type="checkbox"/>	3.mp4 Timestamp - 0 Hours, 00 M...	852.1 KB	application/octet-stream
<input type="checkbox"/>	3.mp4 Timestamp - 0 Hours, 00 M...	849.3 KB	application/octet-stream
<input type="checkbox"/>	3.mp4 Timestamp - 0 Hours, 00 M...	843.4 KB	application/octet-stream
<input type="checkbox"/>	3.mp4 Timestamp - 0 Hours, 00 M...	776.6 KB	application/octet-stream
<input type="checkbox"/>	3.mp4 Timestamp - 0 Hours, 00 M...	780.4 KB	application/octet-stream
<input type="checkbox"/>	3.mp4 Timestamp - 0 Hours, 00 M...	779.8 KB	application/octet-stream
<input type="checkbox"/>	3.mp4 Timestamp - 0 Hours, 00 M...	771.7 KB	application/octet-stream

DETECTION CODE USED

```
for blob in blobs:
    print(blob)
    if blob.name.endswith(".mp4"):
        video_name = os.path.basename(blob.name)
        video_path = f"gs://{new_data_bucket_name}/{blob.name}"
        # Download the video from Google Cloud Storage to the temporary local directory
        temp_video_path = os.path.join(temp_frames_dir, video_name)
        blob.download_to_filename(temp_video_path)
        # Create a folder to store detections for the current video
        video_detections_folder_name = os.path.splitext(video_name)[0]
        video_detections_folder = os.path.join(temp_frames_dir, video_detections_folder_name)
        os.makedirs(video_detections_folder, exist_ok=True)
        # Open the video from the temporary local directory
        cap = cv2.VideoCapture(temp_video_path)
        # Get the video's frames per second (fps) and set the downsampling rate
        fps = cap.get(cv2.CAP_PROP_FPS)
        downsampling_rate = max(int(fps * min_time_difference), 1) seconds
        frame_counter = 0
        # Calculate the fps_multiplier correctly
        fps_multiplier = 1.0 / fps
        while cap.isOpened():
            ret, curr_frame = cap.read()
            if not ret:
                break
            # Only process the frame if it is within the downsampling rate
            if frame_counter % downsampling_rate == 0:
                result = model(curr_frame)
                for res in result:
                    boxes = res.bboxes.cpu().numpy()
                    for i, box in enumerate(boxes):
                        if box.conf[0] > 0.7:
                            res_plotted = res[0].plot(conf=False)
                            current_frame_timestamp = cap.get(cv2.CAP_PROP_POS_MSEC) * fps_multiplier
                            timedelta_obj = str(datetime.timedelta(seconds=round(current_frame_timestamp /
                                1000,1)))


                            x = timedelta_obj.split(':')
                            hh, mm, ss = x[0], x[1], x[2]
                            filename = f"{video_name} Timestamp - {hh} Hours, {mm} Minutes, {ss} Seconds.jpg"
                            output_filename = os.path.join(video_detections_folder, filename)
                            cv2.imwrite(output_filename, res_plotted)
                            print("Saved frame:", output_filename)

            # Increment the frame counter
            frame_counter += 1
        cap.release()
```


RESULTS (DETECTION)



RESULT (SENT TO PUBSUB TOPIC)

 Pub/Sub

Pub/Sub

Topics

Subscriptions

Snapshots

Schemas

Pub/Sub Lite

Lite Reservations

Lite Topics

Lite Subscriptions

Topics

CREATE TOPIC

DELETE

SHOW INFO PANEL

LIST

METRICS

Filter

Filter topics

?

⋮

<input type="checkbox"/>	Topic ID ↑	Encryption key	Topic name	Retention	
<input type="checkbox"/>	animal_detection	Google-managed	projects/gcds-oht33765u9-2023/topics/animal_detection	—	⋮

SUGGESTED CHANGE 1 – AUTOMATING DETECTION RUN THROUGH SERVERLESS CLOUD FUNCTIONS



Videos from Tap Camera come in batches daily.

e.g. 24 1-hour videos for yesterday

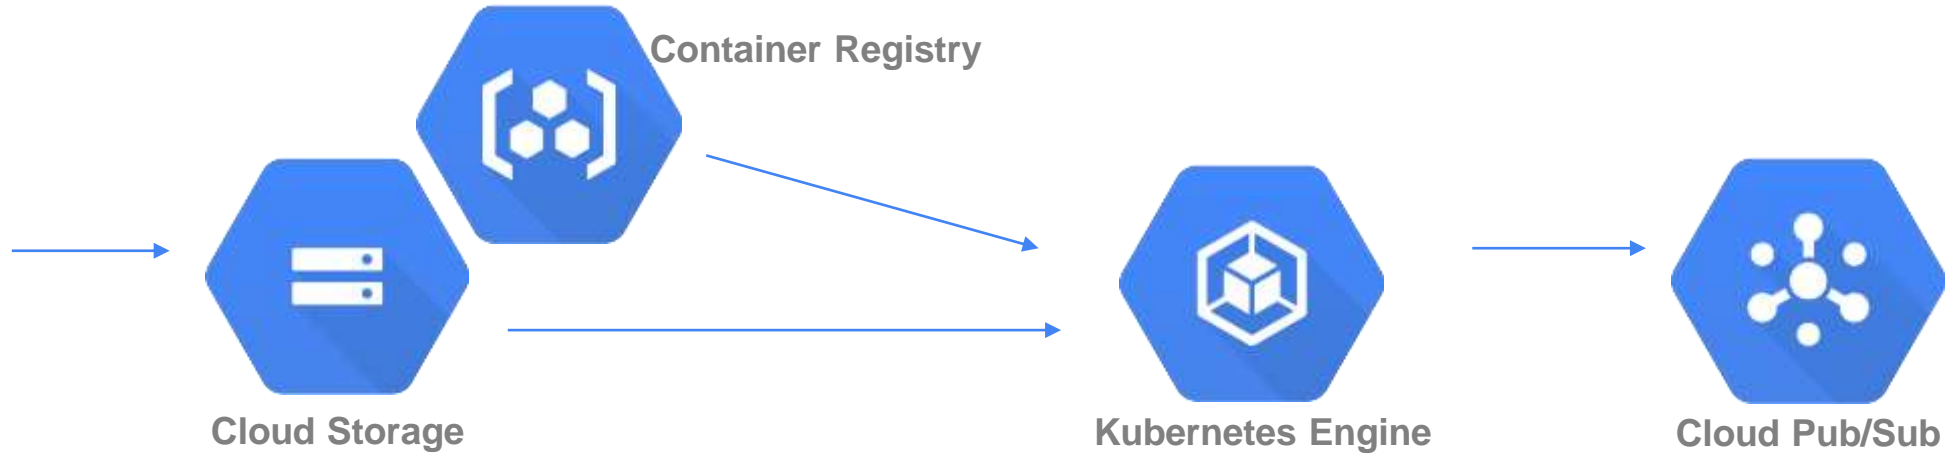
The videos land in a particular folder inside a bucket in Cloud Storage

As soon as a new video lands in the bucket, using object finalization trigger, Cloud Functions can kick in and run a python code to pass the videos through an object detection model.

The code automatically stores the detections in the bucket. The detections are images with bounding boxes around the animal found. Also, the detections are named as VideoName-Timestamp.jpg

As soon as new detections are stored in the bucket, Cloud PubSub will trigger a notification to subscribed web services, which can ingest the topic to send email/SMS to concerned people

SUGGESTED CHANGE 2 — AUTOMATING DETECTION RUN THROUGH KUBERNETES CONTAINERIZATION



Videos from Tap Camera come in batches daily.

e.g. 24 1-hour videos for yesterday

The videos land in a particular folder inside a bucket in Cloud Storage

The Kubernetes Engine can host the model as an endpoint on a webservice, using which the videos can be processed upon.

The model along with all the dependencies will be present in the container registry as a docker container.

Then the endpoint can send the detections directly to pubsub.

As soon as new detections are stored in the bucket, Cloud PubSub will trigger a notification to subscribed web services, which can ingest the topic to send email/SMS to concerned people

DRONE SEARCH & RESCUE

UseCase Type:

Model Deployment : Edge Device (Drone Camera)

Detection : Edge Device

Storage : Cloud

Action Taken : On Cloud (Alerting Concerned Services/ People)

OVERVIEW



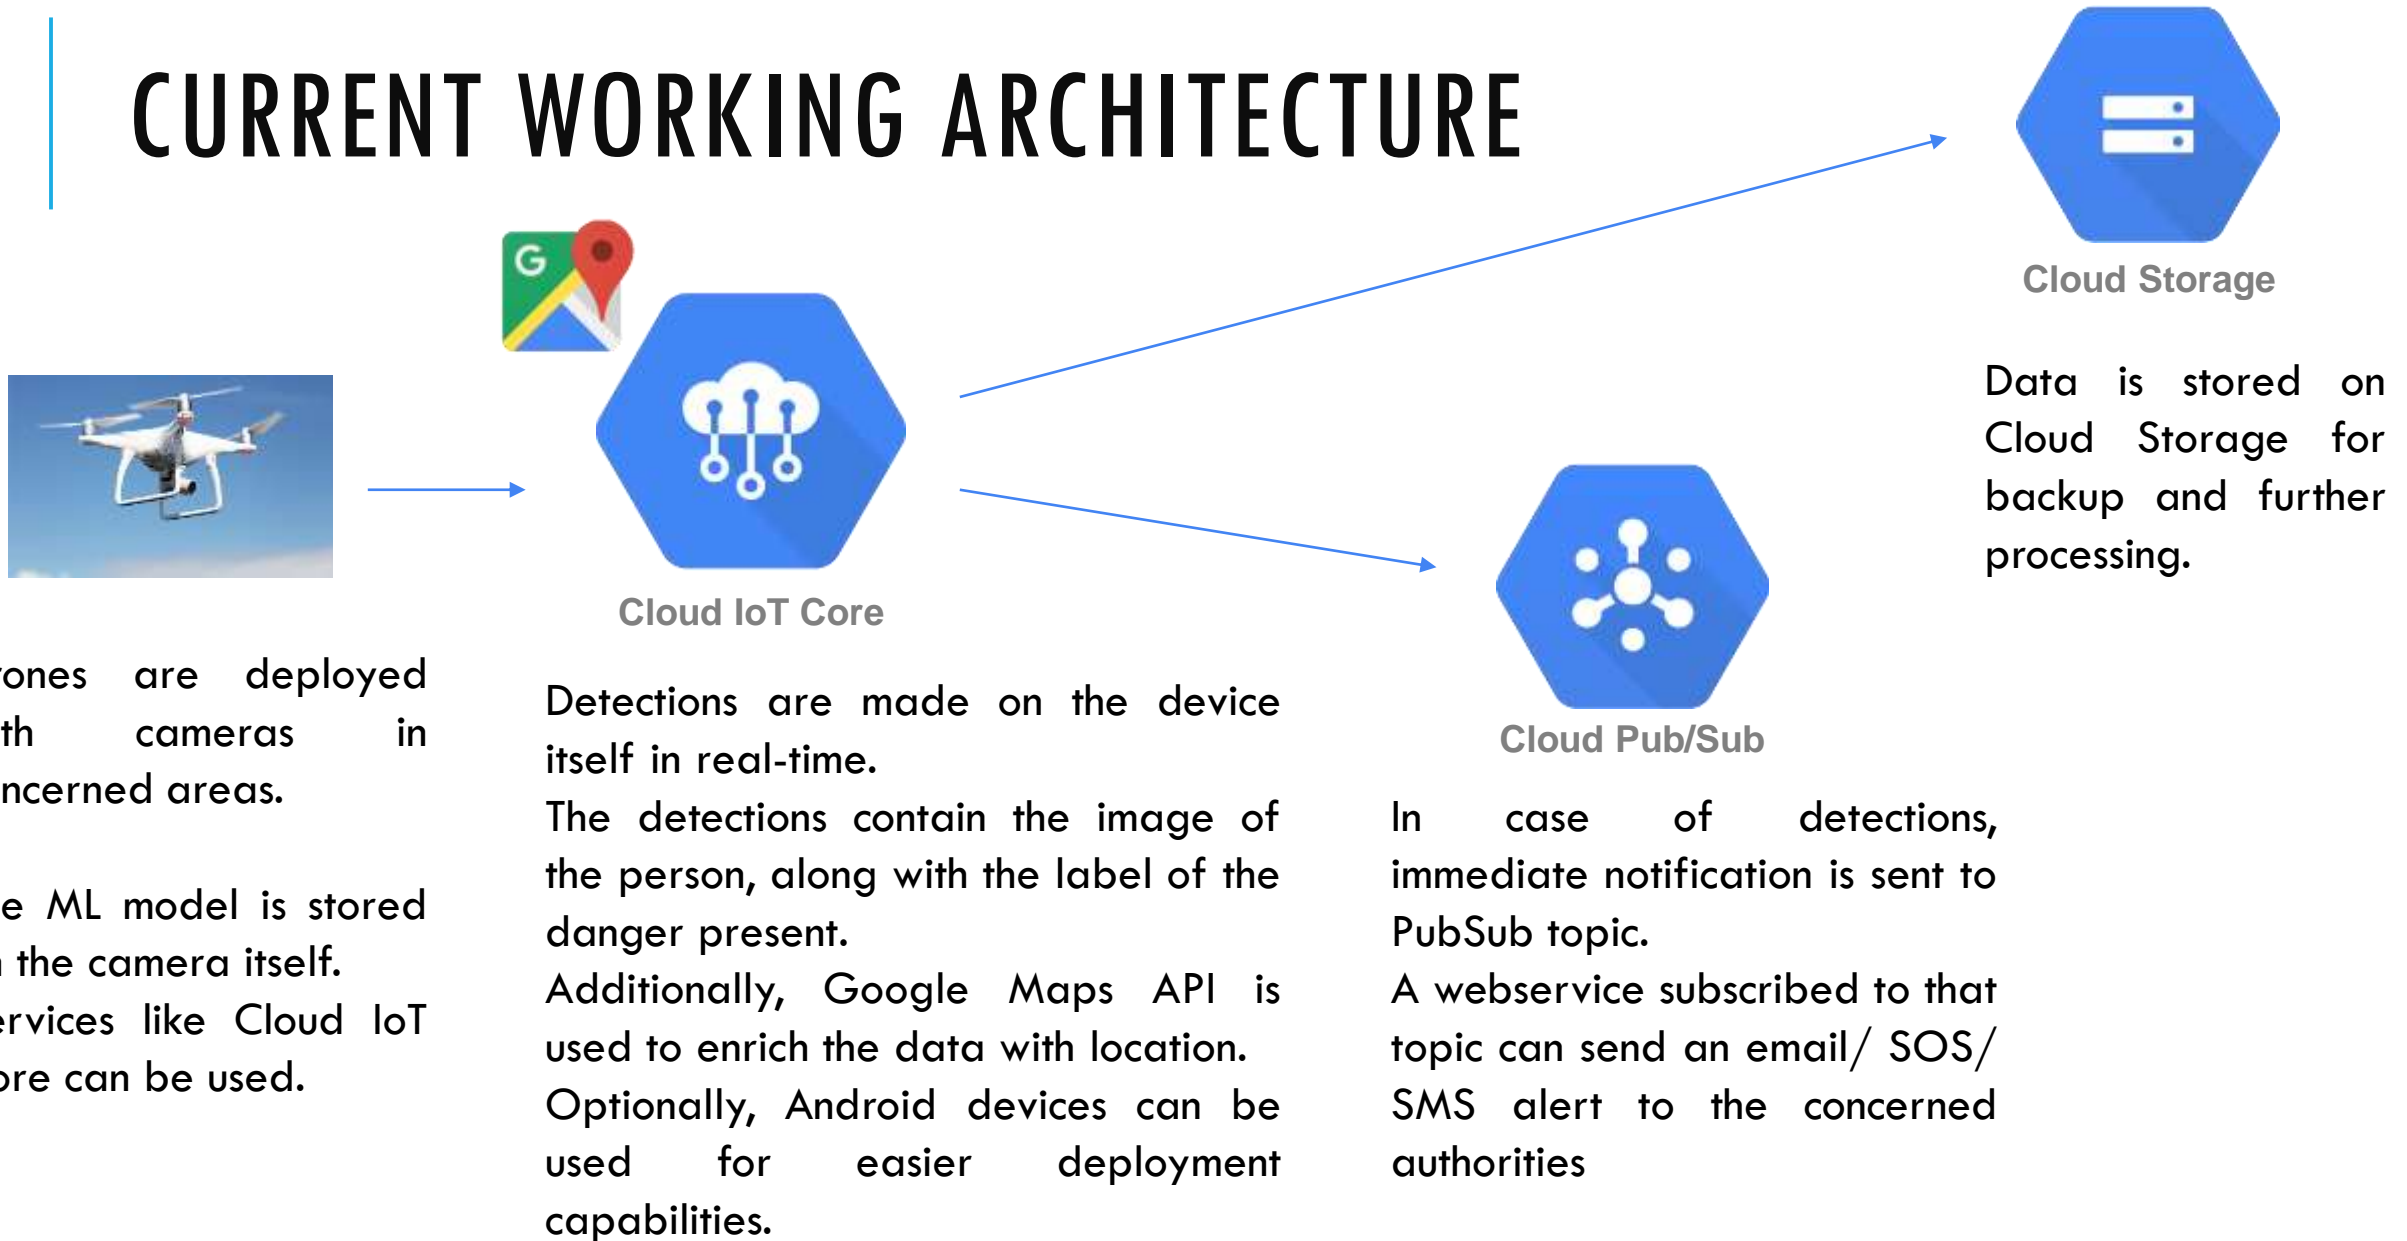
Drones can be used to detect people and animals requiring urgent help from immediate physical harm in remote areas or hard-to-monitor places.

As soon as a detection is made, the concerned authorities can be notified and alerted to take immediate actions.

Moreover, an accurate location can be sent along with the detection.



CURRENT WORKING ARCHITECTURE



Drones are deployed with cameras in concerned areas.

The ML model is stored on the camera itself. Services like Cloud IoT Core can be used.

Detections are made on the device itself in real-time.

The detections contain the image of the person, along with the label of the danger present.

Additionally, Google Maps API is used to enrich the data with location. Optionally, Android devices can be used for easier deployment capabilities.

In case of detections, immediate notification is sent to PubSub topic.

A webservice subscribed to that topic can send an email/ SOS/ SMS alert to the concerned authorities

EXAMPLE SCENARIO



In our beach park, we would like to keep track of majorly two types of accidents:

1. Drowning People,
2. Capsized Boats


Since the life guards are not able to monitor far away waters, we would like to take the help of a drone continuously monitoring the waters.

The drone should alert us immediately in case of the above two dangers, following which we can send a rapid response team.

The detection should contain an image of the human(s), with the danger, and also the location of the event.

MODEL USED

YOLOv8 Object Detection Model (<https://docs.ultralytics.com/>)



MODEL TRAINING

The model was trained on Vertex AI Notebook & Google Colab, with 800+ images collected and labelled manually across the 2 classes in the YOLO format.

The dataset was further augmented to include multiple scenarios for light, color, distortions etc.

```
from ultralytics import YOLO
import locale
locale.getpreferredencoding = lambda: "UTF-8"
from google.colab.patches import cv2_imshow
import tensorflow as tf
#Avoid OOM errors by setting GPU Memory Consumption Growth
gpus = tf.config.experimental.list_physical_devices('GPU')
for gpu in gpus:
    tf.config.experimental.set_memory_growth(gpu, True)
tf.config.list_physical_devices('GPU')
from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

model = YOLO("yolov8m.pt")






Downloading https://github.com/ultralytics/assets/releases/download/v0.0.0/yolov8m.pt to yolov8m.pt...
100% [49.7M/49.7M [00:00<00:00, 56.7MB/s]]

model.train(data=r"/content/drive/MyDrive/Machine Learning/Drone Search and Rescue YOLO/dataset/dataset.yaml", epochs=100, project=
```

UPLOAD FILES UPLOAD FOLDER CREATE FOLDER

Filter by name prefix only ▾

≡ Filter Filter objects and folders

<input type="checkbox"/>	Name	Size
<input type="checkbox"/>	 bestn.pt ← Final Model	5.9 MB
<input type="checkbox"/>	 dataset_0/	—
<input type="checkbox"/>	 logs/	—
<input type="checkbox"/>	 process_video.py	3.7 KB
<input type="checkbox"/>	 runs_0/	—

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
99/100	7.25G	0.8303	0.6203	1.239	3	640: 100% [██████████] 11/11 [00:05<00:00, 1.99it/s]
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% [██████████] 2/2 [00:00<00:00, 2.59it/s]
	all	41	41	0.884	0.685	0.816 0.391

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
100/100	7.22G	0.8804	0.6307	1.301	3	640: 100% [██████████] 11/11 [00:05<00:00, 2.00it/s]
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% [██████████] 2/2 [00:02<00:00, 1.19s/it]
	all	41	41	0.877	0.685	0.82 0.388

100 epochs completed in 0.292 hours.

Optimizer stripped from /content/drive/MyDrive/Machine Learning/Drone Search and Rescue YOLO/dataset/train3/weights/last.pt, 52.0MB

Optimizer stripped from /content/drive/MyDrive/Machine Learning/Drone Search and Rescue YOLO/dataset/train3/weights/best.pt, 52.0MB

Validating /content/drive/MyDrive/Machine Learning/Drone Search and Rescue YOLO/dataset/train3/weights/best.pt...

Ultralytics YOLOv8.0.132 Python-3.10.12 torch-2.0.1+cu118 CUDA:0 (Tesla T4, 15102MiB)

Model summary (fused): 218 layers, 25840918 parameters, 0 gradients

Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% [██████████] 2/2 [00:01<00:00, 1.87it/s]
all	41	41	0.948	0.697	0.802	0.431
capsized_boat	41	35	0.896	0.737	0.909	0.448
drowning_person	41	6	1	0.656	0.696	0.414

Speed: 5.9ms preprocess, 8.9ms inference, 0.0ms loss, 1.5ms postprocess per image

Results saved to /content/drive/MyDrive/Machine Learning/Drone Search and Rescue YOLO/dataset/train3

SAMPLE OPERATION — STEP 1

As an example, 2 videos were placed in a specific folder:

[←](#) Bucket details [REFRESH](#) [HELP ASSISTANT](#) [LEARN](#)

drone-search-rescue-data

Location	Storage class	Public access	Protection
europa-west6 (Zurich)	Standard	Not public	None

[OBJECTS](#) [CONFIGURATION](#) [PERMISSIONS](#) [PROTECTION](#) [LIFECYCLE](#) [OBSERVABILITY](#) [INVENTORY REPORTS](#)

Buckets > drone-search-rescue-data > new_data

[UPLOAD FILES](#) [UPLOAD FOLDER](#) [CREATE FOLDER](#) [TRANSFER DATA](#) [MANAGE HOLDS](#) [DOWNLOAD](#) [DELETE](#)

Filter by name prefix only

Filter | Filter objects and folders

Show deleted data

<input type="checkbox"/>	Name	Size	Type	Created	Storage class	Last modified	Public access	
<input type="checkbox"/>	1.mp4	9.6 MB	video/mp4	Jul 20, 2023, 3:10:34 PM	Standard	Jul 20, 2023, 3:10:34 PM	Not public	
<input type="checkbox"/>	2.mp4	4.5 MB	video/mp4	Jul 20, 2023, 3:10:33 PM	Standard	Jul 20, 2023, 3:10:33 PM	Not public	

SAMPLE OPERATION – STEP 2

After the model was run in the Vertex AI notebook, it created the detection images in separate folder for each video, with the image name as VideoName-Timestamp.jpg

Buckets > drone-search-rescue-data > detections

UPLOAD FILES UPLOAD FOLDER CREATE FOLDER TRANSFER

Filter by name prefix only ▾ Filter Filter objects and folders

<input type="checkbox"/>	Name	Size	Type	Created
<input type="checkbox"/>	1/	—	Folder	—
<input type="checkbox"/>	2/	—	Folder	—

Buckets > drone-search-rescue-data > detections > 2

UPLOAD FILES UPLOAD FOLDER CREATE FOLDER TRANSFER DATA ▾ MANAGE HOLDS DOWNLOAD DELETE

Filter by name prefix only ▾ Filter Filter objects and folders Show deleted data

<input type="checkbox"/>	Name	Size	Type	Created	Storage class	Last modified	P
<input type="checkbox"/>	2.mp4 Timestamp - 0 Hours, 00 M...	70.2 KB	application/octet-stream	Jul 20, 2023, 5:02:44 PM	Standard	Jul 20, 2023, 5:02:44 PM	N
<input type="checkbox"/>	2.mp4 Timestamp - 0 Hours, 00 M...	69.8 KB	application/octet-stream	Jul 20, 2023, 5:02:43 PM	Standard	Jul 20, 2023, 5:02:43 PM	N
<input type="checkbox"/>	2.mp4 Timestamp - 0 Hours, 00 M...	76.8 KB	application/octet-stream	Jul 20, 2023, 5:02:43 PM	Standard	Jul 20, 2023, 5:02:43 PM	N
<input type="checkbox"/>	2.mp4 Timestamp - 0 Hours, 00 M...	91 KB	application/octet-stream	Jul 20, 2023, 5:02:44 PM	Standard	Jul 20, 2023, 5:02:44 PM	N
<input type="checkbox"/>	2.mp4 Timestamp - 0 Hours, 00 M...	80.9 KB	application/octet-stream	Jul 20, 2023, 5:02:43 PM	Standard	Jul 20, 2023, 5:02:43 PM	N
<input type="checkbox"/>	2.mp4 Timestamp - 0 Hours, 00 M...	39.6 KB	application/octet-stream	Jul 20, 2023, 5:02:44 PM	Standard	Jul 20, 2023, 5:02:44 PM	N
<input type="checkbox"/>	2.mp4 Timestamp - 0 Hours, 00 M...	47 KB	application/octet-stream	Jul 20, 2023, 5:02:44 PM	Standard	Jul 20, 2023, 5:02:44 PM	N
<input type="checkbox"/>	2.mp4 Timestamp - 0 Hours, 00 M...	47.3 KB	application/octet-stream	Jul 20, 2023, 5:02:46 PM	Standard	Jul 20, 2023, 5:02:46 PM	N

DETECTION CODE USED

```
for blob in blobs:
    print(blob)
    if blob.name.endswith(".mp4"):
        video_name = os.path.basename(blob.name)
        video_path = f"gs://{new_data_bucket_name}/{blob.name}"
        # Download the video from Google Cloud Storage to the temporary local directory
        temp_video_path = os.path.join(temp_frames_dir, video_name)
        blob.download_to_filename(temp_video_path)
        # Create a folder to store detections for the current video
        video_detections_folder_name = os.path.splitext(video_name)[0]
        video_detections_folder = os.path.join(temp_frames_dir, video_detections_folder_name)
        os.makedirs(video_detections_folder, exist_ok=True)
        # Open the video from the temporary local directory
        cap = cv2.VideoCapture(temp_video_path)
        # Get the video's frames per second (fps) and set the downsampling rate
        fps = cap.get(cv2.CAP_PROP_FPS)
        downsampling_rate = max(int(fps * min_time_difference), 1) seconds
        frame_counter = 0
        # Calculate the fps_multiplier correctly
        fps_multiplier = 1.0 / fps
        while cap.isOpened():
            ret, curr_frame = cap.read()
            if not ret:
                break
            # Only process the frame if it is within the downsampling rate
            if frame_counter % downsampling_rate == 0:
                result = model(curr_frame)
                for res in result:
                    boxes = res.bboxes.cpu().numpy()
                    for i, box in enumerate(boxes):
                        if box.conf[0] > 0.7:
                            res_plotted = res[0].plot(conf=False)
                            current_frame_timestamp = cap.get(cv2.CAP_PROP_POS_MSEC) * fps_multiplier
                            timedelta_obj = str(datetime.timedelta(seconds=round(current_frame_timestamp /
                                1000,1)))

                            x = timedelta_obj.split(':')
                            hh, mm, ss = x[0], x[1], x[2]
                            filename = f"{video_name} Timestamp - {hh} Hours, {mm} Minutes, {ss} Seconds.jpg"
                            output_filename = os.path.join(video_detections_folder, filename)
                            cv2.imwrite(output_filename, res_plotted)
                            print("Saved frame:", output_filename)

            # Increment the frame counter
            frame_counter += 1
        cap.release()
```

RESULTS (DETECTION)



SUGGESTED CHANGES

- Training the model with better data i.e. top view data taken from a drone, with classes simulated in real time.
- Training for multiple classes i.e. other dangers e.g. sharks
- Training for multiple angles i.e. looking down from drone, looking straight etc.
- Deployment of Android Devices on the Drone using Raspberry Pi, instead of expensive drone cameras.
- Drones can also carry first aid kit and small supplies, that can be dropped to people who need them in the interim the rescue people come.

DRIVER DROWSINESS DETECTION

UseCase Type:

Model Deployment : Edge Device (Vehicle Dash Camera)

Detection : Edge Device

Storage : Edge Device, Cloud

Action Taken : On Device (Raising Car Alarm, Stopping Car)

OVERVIEW



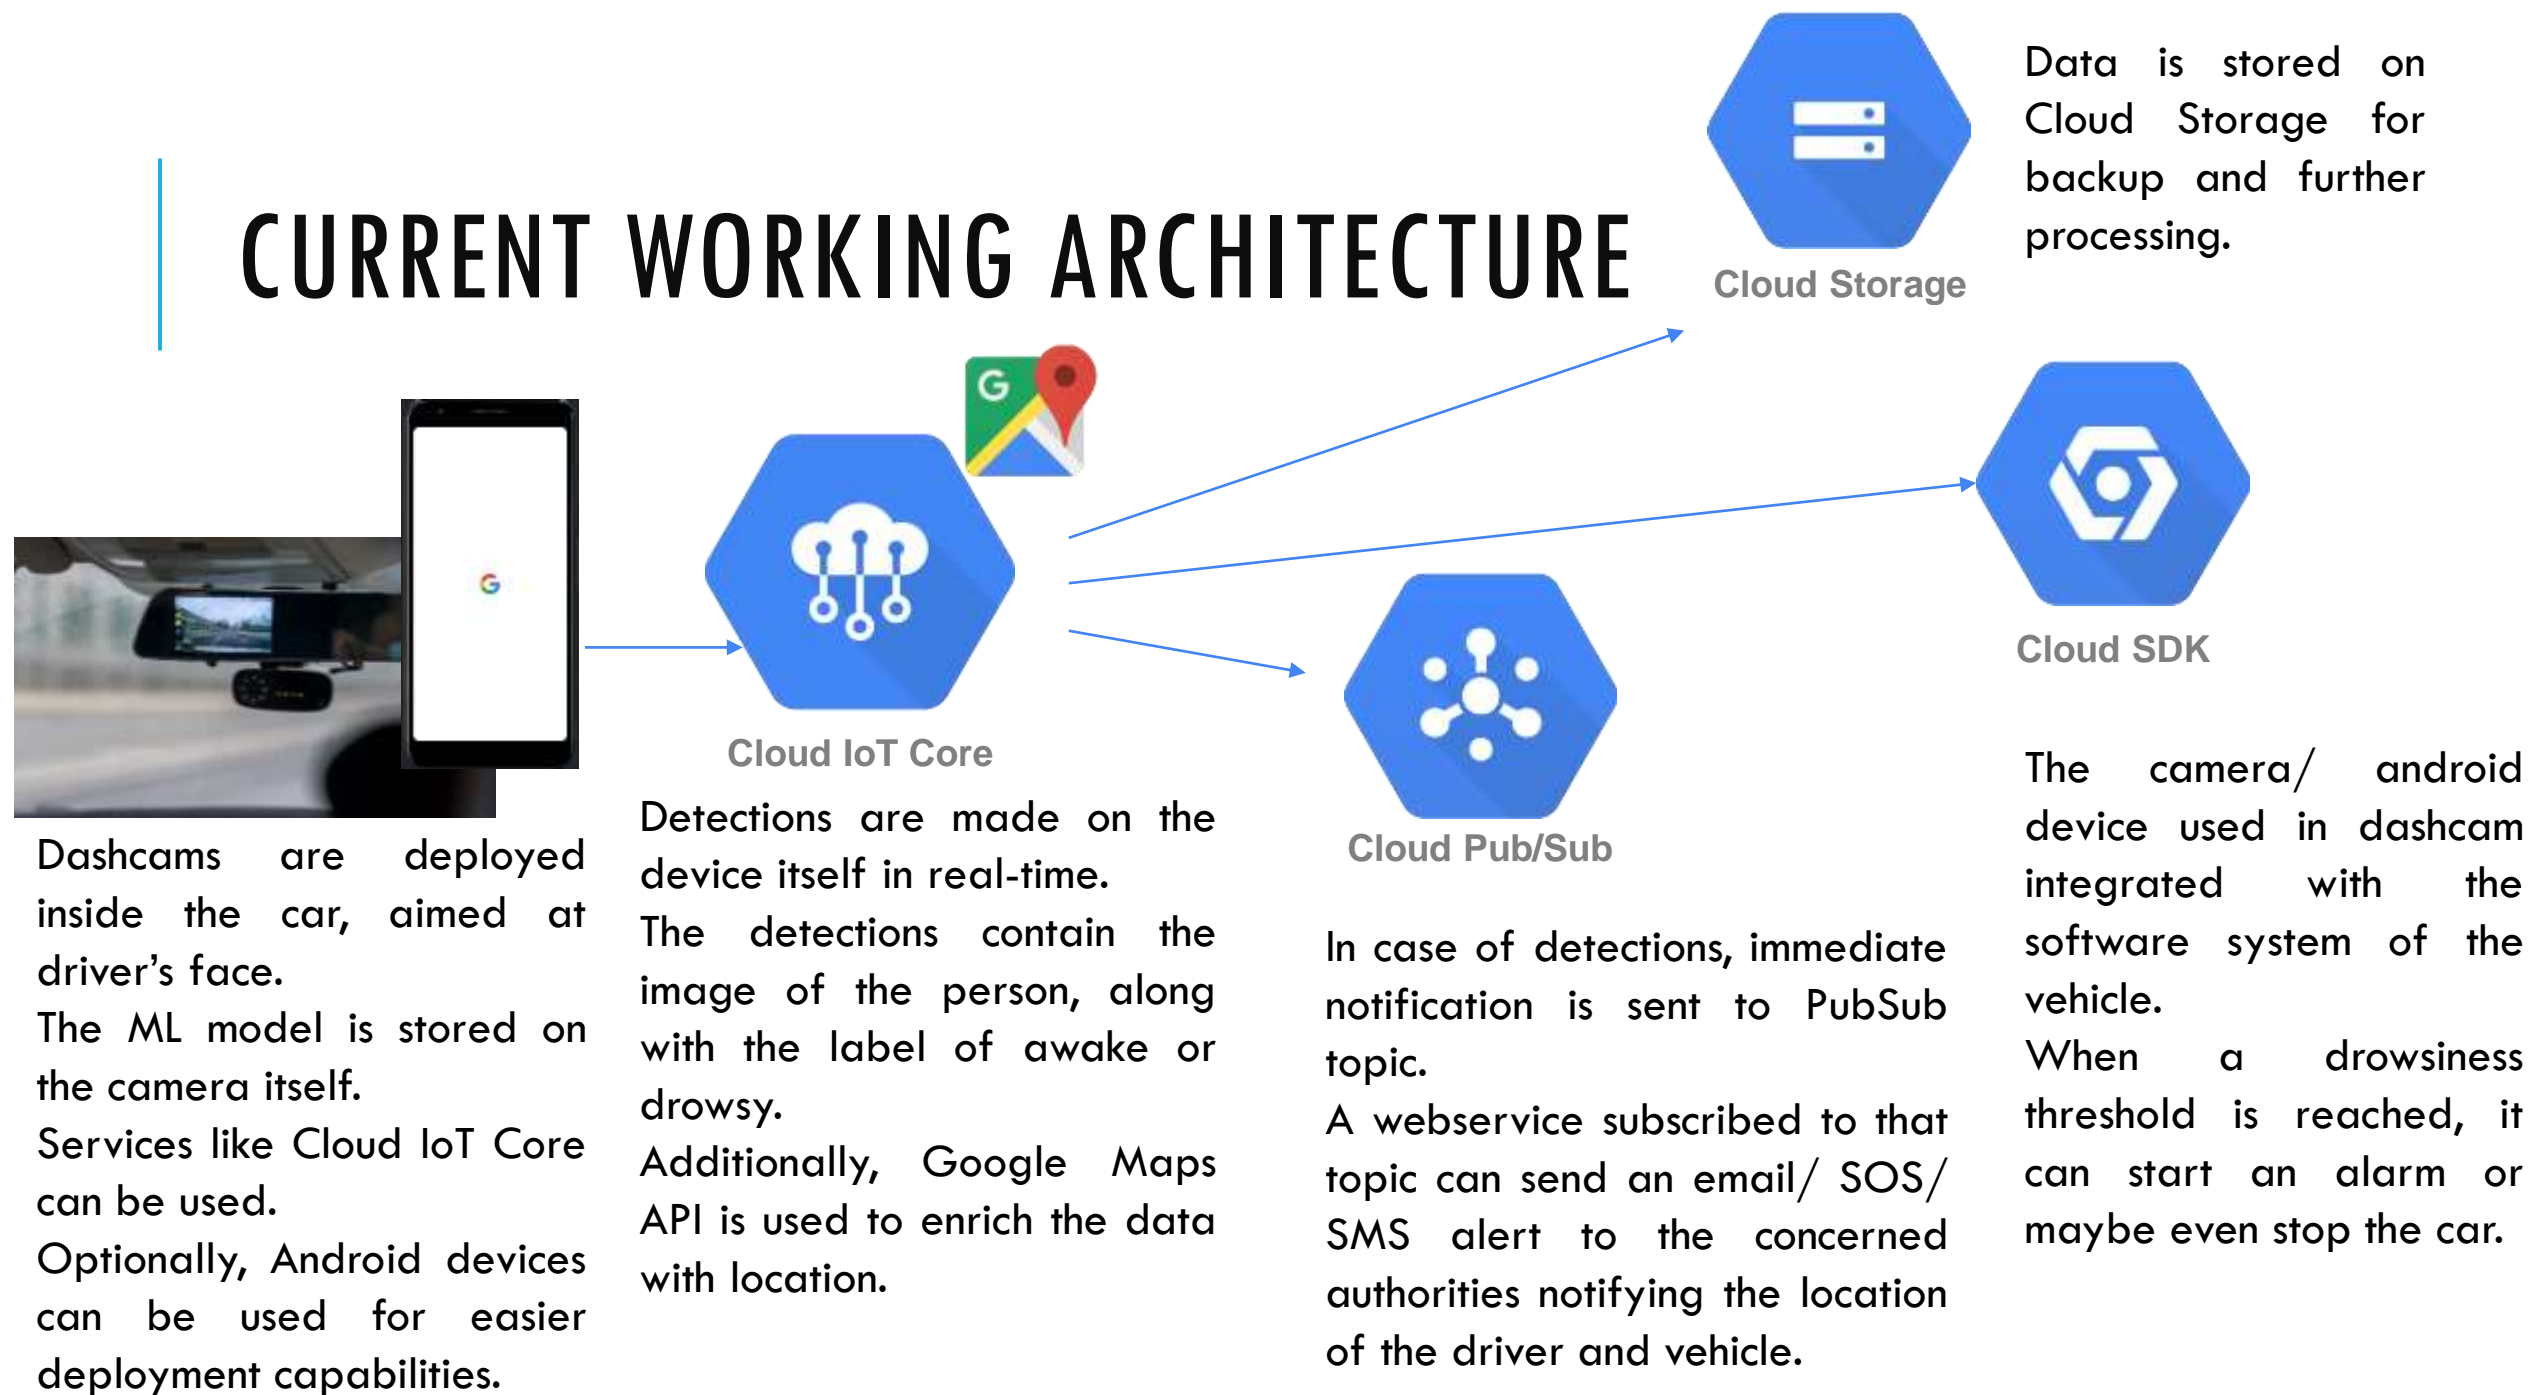
Dashboard cameras (Dashcams) are getting increasingly popular. They can be both inward facing (looking inside the car), as well as outward facing (looking outside the car, typically in front of the car).

They are used to monitor and store the actions of passengers, drivers, pedestrians, traffic etc by the fleet owners of cabs, trucks etc.

A major concern for all the parties involved is the drowsiness of the driver, which can cause fatal accidents.

Though the drowsiness cannot be prevented, measures can be taken such that as soon as a driver feels drowsy, he can be awoken by an instant alarm in the vehicle, or maybe even the vehicle can be brought to a halt safely.

CURRENT WORKING ARCHITECTURE




EXAMPLE SCENARIO

- We run a fleet of shipping trucks.
- We need to constantly monitor our drivers for signs of drowsiness.
- If a driver is drowsy for more than 5 seconds, an alarm should ring in the car.
- The alarm volume should keep on rising if the driver is still drowsy.
- If the said driver is drowsy for more than 10 seconds, stop the vehicle.
- Also send us a notification with a location, in case a vehicle is stopped



MODEL USED

VGG16 Classification Model (<https://keras.io/api/applications/vgg/>)



MODEL TRAINING

The model was trained on local machine, with 4000+ images collected and labelled across the 2 classes (drowsy & awake) in a Kaggle repository.

(<https://www.kaggle.com/datasets/prasadvp/atil/mrl-dataset>)

The repository contains only the data for eyes, which is a good starting point for training a drowsiness model, as drowsiness corresponds to longer periods of closed eyes.

The dataset was further augmented to include multiple scenarios for light, color, distortions etc.

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout

model = Sequential()

#
model.add(Conv2D(16, (3,3), 1, activation='relu', input_shape=(256,256,3)))
model.add(MaxPooling2D())

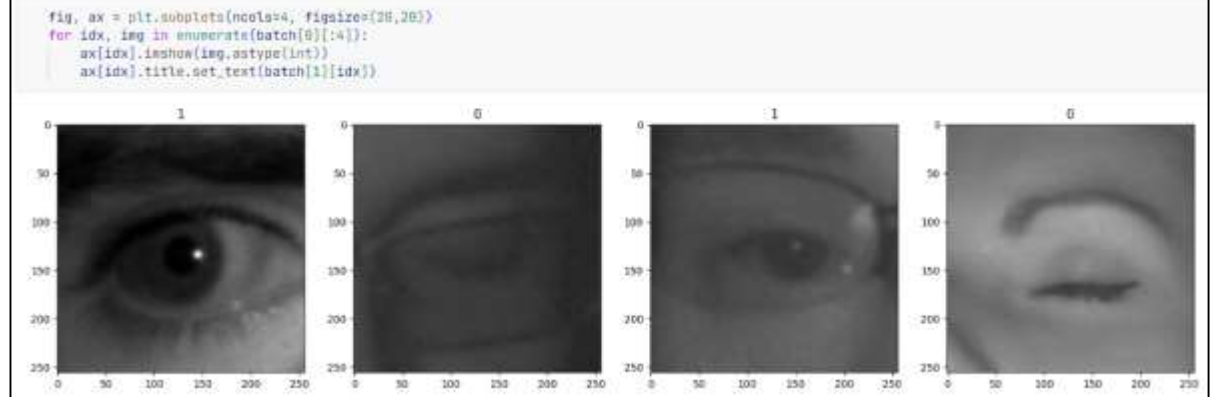
model.add(Conv2D(32, (3,3), 1, activation='relu'))
model.add(MaxPooling2D())

model.add(Conv2D(16, (3,3), 1, activation='relu'))
model.add(MaxPooling2D())

model.add(Flatten())

model.add(Dense(256, activation='relu'))
model.add(Dense(1, activation='sigmoid'))

model.compile('adam', loss=tf.losses.BinaryCrossentropy(), metrics=['accuracy'])
```



```
new_model = load_model(os.path.join('models', 'cnnimageclassifier.h5'))

yhatnew = new_model.predict(np.expand_dims(resize/255, 0))

1/1 [*****] - 0s 138ms/step

if yhatnew > 0.5:
    print(f'Predicted class is Awake')
else:
    print(f'Predicted class is Drowsy')
```


SAMPLE OPERATION

As an example, 1 video was placed in a specific folder.

The video shows a man driving a truck while being drowsy.

The video was passed through the model as shown in the code.

```
from tensorflow.keras.models import load_model
new_model =
load_model(r'C:\Users\HP\Desktop\drowsinessdetection.h5')
def detect_objects_in_video(model, video_path):
    # Open the video file
    cap = cv2.VideoCapture(video_path)
    while True:
        ret, frame = cap.read()
        if not ret:
            break
        # Preprocess the frame (adjust according to your model's
input requirements)
        frame = cv2.cvtColor(frame, cv2.COLOR_BGR2RGB)
        frame = cv2.resize(frame, (256, 256)) # Adjust input
dimensions
        frame = frame / 255.0 # Normalize pixel values
        # Make predictions using the model
        prediction = model.predict(np.expand_dims(frame, axis=0))
        if prediction > 0.5:
            print(f'Predicted class is Awake')
        else:
            print(f'Predicted class is Drowsy')

        # Process the predictions and draw bounding boxes
        # (You need to customize this part based on your model's
output and the objects you want to detect)

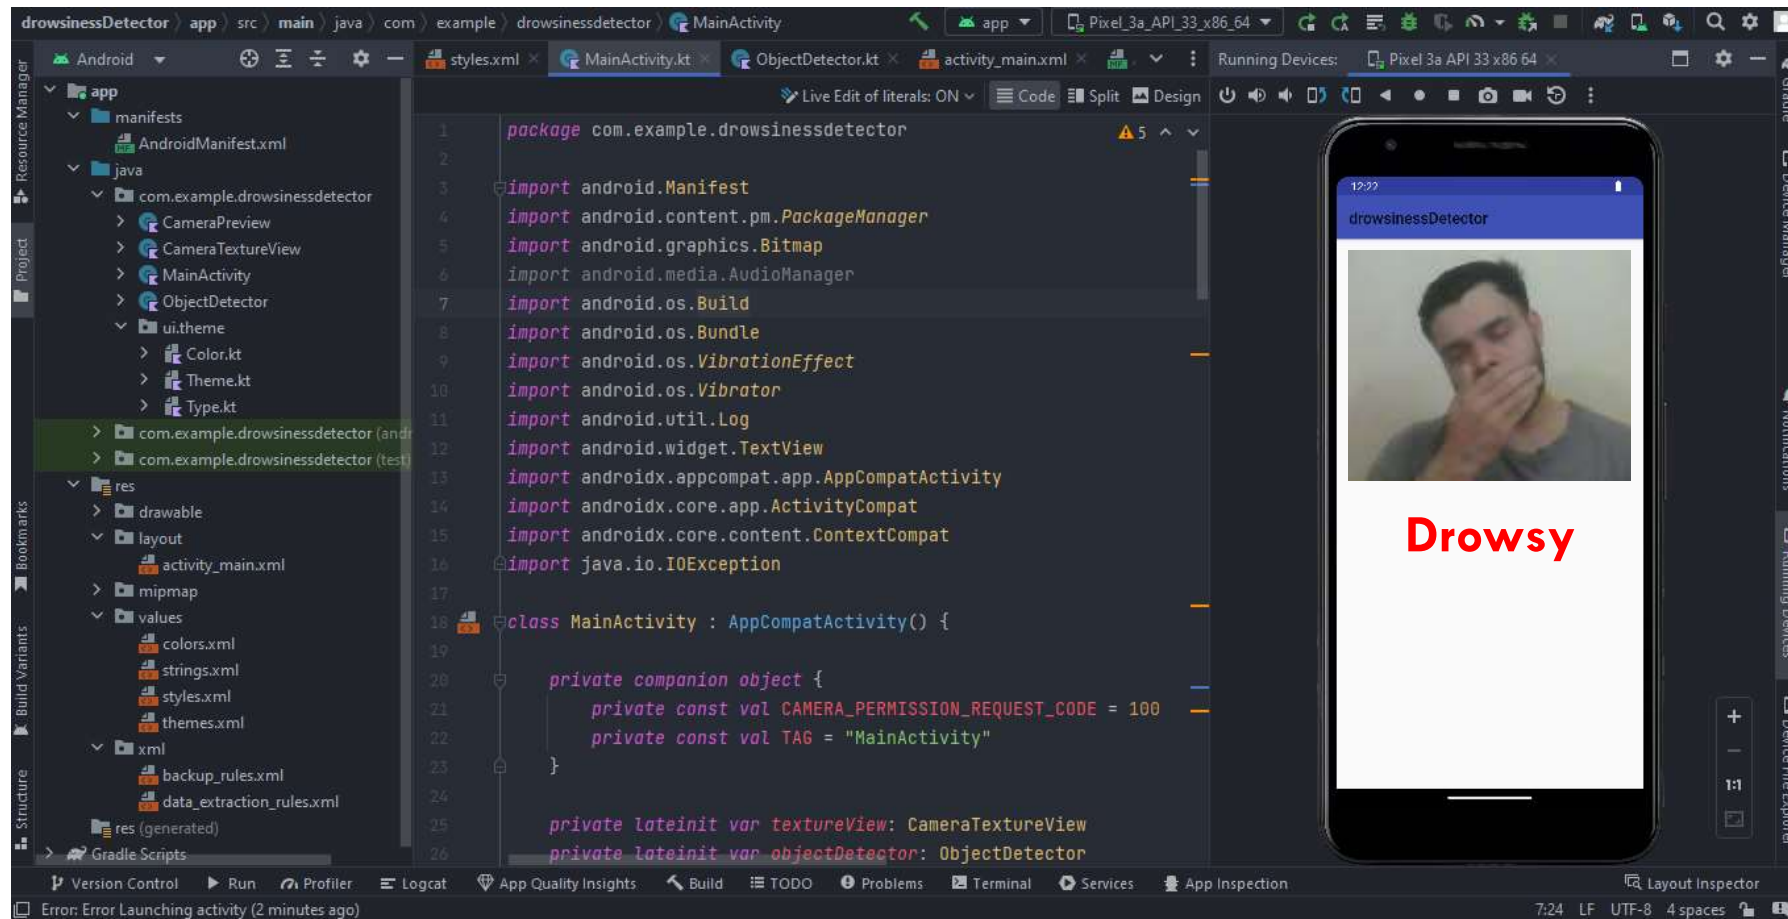
        # Display the frame with detections
        cv2.imshow('Object Detection', frame)
        # Exit on pressing 'q'
        if cv2.waitKey(1) & 0xFF == ord('q'):
            break
    # Release video capture and close the OpenCV window
    cap.release()
    cv2.destroyAllWindows()
# Replace 'video_path' with the path to your video file
video_path = r'C:\Users\HP\Desktop\1.mp4'
detect_objects_in_video(new_model, video_path)
```



Drowsy

Awake

ANDROID APP AS A SUBSTITUTE FOR DASHCAM



The Android App takes in the video feed from the front camera, and applies to it a TFLite version of the model stored in the application package.

The model is run continuously on each of the frames, and classifies it as drowsy or awake.

If the drowsy class is maintained for more than 5 seconds, the phone vibrates.

SUGGESTED CHANGES

- Refined Android App
- An ML model trained on an exhaustive set of data for drowsiness, which includes yawning, hand movement, etc.
- Integration of the Android App with vehicle's electrical system.
- Deployment of the package on interior facing dashcams focused on the driver's face.